



NATIONAL UNIVERSITY OF SCIENCE AND TECHNOLOGY

DEPARTMENT OF COMPUTER SCIENCE

ARTIFICIAL INTELLIGENCE LAB

NAME	Ayesha Imran
Class	CS-A
Lab	01
Course	Artificial Intelligence
Date	1-October-25
Submitted To	Lec. Ijlal Haider

IN LAB TASKS

Task 1:

Implement a Simple Reflex Agent for a Grid-Based Environment with the following Requirements:

1. The agent should suck the dirt if the current square is dirty. This is the primary action
the agent performs when it detects dirt in its current square.
2. If the square is clean, the agent should move randomly in one of the available directions (north, south, west, east). However, the agent should avoid moving into walls (obstacles). The agent is provided with bumper sensor data indicating whether there is a wall in each direction.
3. The agent cannot directly access any variables in the environment. Instead, it only receives:
 - Bumper Sensor Data: A dictionary that indicates whether there is a wall (bumper) in each of the four directions (north, south, west, east).
 - Dirt Sensor: A Boolean value that indicates whether the current square is dirty.
4. The simulation stops when either:
 - All squares in the grid are cleaned, or
 - The agent reaches the maximum number of steps.

Solution:

```
import random

class SimpleReflexAgent:
    def __init__(self, grid, max_steps=20):
        self.grid = grid
        self.rows = len(grid)
        self.cols = len(grid[0])
        self.position = (0, 0)      # agent starts at top-left
        self.max_steps = max_steps
        self.steps = 0

    # Check if current cell is dirty
    def dirt_sensor(self):
        r, c = self.position
        return self.grid[r][c] == 1

    # Check where the walls are
    def bumper_sensor(self):
```

```

        r, c = self.position
        return {
            "north": r == 0,
            "south": r == self.rows - 1,
            "west": c == 0,
            "east": c == self.cols - 1
        }

# Clean the current cell
def suck(self):
    r, c = self.position
    self.grid[r][c] = 0
    print(f"Step {self.steps}: Cleaned at {self.position}")

# Move randomly to a free direction
def move(self):
    walls = self.bumper_sensor()
    directions = [d for d, is_wall in walls.items() if not is_wall] # only
valid moves

    direction = random.choice(directions)
    r, c = self.position

    if direction == "north": r -= 1
    elif direction == "south": r += 1
    elif direction == "west": c -= 1
    elif direction == "east": c += 1

    self.position = (r, c)
    print(f"Step {self.steps}: Moved {direction} to {self.position}")

# Check if all cells are clean
def all_clean(self):
    for row in self.grid:
        if 1 in row:
            return False
    return True

# Run the agent
def run(self):
    while self.steps < self.max_steps:
        self.steps += 1

        if self.dirt_sensor():
            self.suck()

```

```

        else:
            self.move()

            if self.all_clean():
                print("▽ All clean! Done.")
                return
            print("▣ Max steps reached. Stopping.")

# -----
# MAIN PROGRAM
# -----
grid = [
    [1, 0, 1],
    [1, 1, 0],
    [0, 1, 1]
]

agent = SimpleReflexAgent(grid, max_steps=20)
agent.run()

```

Output:

Step 1: Cleaned at (0, 0)

Step 2: Moved east to (0, 1)

Step 3: Moved south to (1, 1)

Step 4: Cleaned at (1, 1)

Step 5: Moved east to (1, 2)

Step 6: Moved west to (1, 1)

Step 7: Moved south to (2, 1)

Step 8: Cleaned at (2, 1)

Step 9: Moved north to (1, 1)

Step 10: Moved south to (2, 1)

Step 11: Moved north to (1, 1)

Step 12: Moved west to (1, 0)

Step 13: Cleaned at (1, 0)

Step 14: Moved north to (0, 0)

Step 15: Moved south to (1, 0)

Step 16: Moved east to (1, 1)

Step 17: Moved east to (1, 2)

Step 18: Moved north to (0, 2)

Step 19: Cleaned at (0, 2)

Step 20: Moved west to (0, 1)

Max steps reached. Stopping.

Task 2:

Implement a Model-Based Reflex Agent for a Grid-Based Environment. The agent's decisions are based on both its percepts (sensor readings) and its internal state. The agent needs to track its environment's cleanliness and avoid revisiting locations that it has already cleaned.

1. If the agent detects that the current square is dirty, it should clean the square.
2. The agent should move to an adjacent square (north, south, west, or east) if there are no bumpers (walls) in that direction and the square has not already been visited.
3. The agent should maintain an internal map of cleanliness and a record of visited positions to avoid revisiting and cleaning the same square more than once.
4. The agent has access to bumper sensor data which tells it if there are walls in each of the four directions (north, south, west, east).
5. The agent will move and clean the environment until:
 - All squares are cleaned,
 - The agent reaches the maximum number of steps allowed.

Solution:

```
import numpy as np

class ModelBasedReflexAgent:
    def __init__(self, grid_shape):
        self.visited = set()
        self.grid_shape = grid_shape

    def next_action(self, pos, bumpers, dirty):
        self.visited.add(tuple(pos))
        if dirty:
            return "suck"
        # Find unvisited, reachable neighbors
        moves = []
        directions = {
            "north": (-1, 0),
            "south": (1, 0),
            "west": (0, -1),
            "east": (0, 1)
        }
        for d, (dr, dc) in directions.items():
            if not bumpers[d]:
                new_r, new_c = pos[0] + dr, pos[1] + dc
                if (0 <= new_r < self.grid_shape[0] and
                    0 <= new_c < self.grid_shape[1] and
                    (new_r, new_c) not in self.visited):
                    moves.append(d)
        if moves:
            return np.random.choice(moves)
        # If all neighbors visited, move randomly to any open direction
        open_moves = [d for d in directions if not bumpers[d]]
        return np.random.choice(open_moves) if open_moves else None

def environment_model_based(agent, grid, max_steps=50):
    pos = [0, 0]
    cleaned = 0
    steps = 0
    while steps < max_steps and np.any(grid == 1):
        r, c = pos
        dirty = grid[r, c] == 1
        bumpers = {
            "north": r == 0,
            "south": r == grid.shape[0] - 1,
            "west": c == 0,
```

```

        "east": c == grid.shape[1] - 1
    }
    action = agent.next_action(pos, bumpers, dirty)
    print(f"Step {steps+1}: Pos {pos}, Dirty: {dirty}, Action: {action}")
    if action == "suck":
        grid[r, c] = 0
        cleaned += 1
    elif action == "north" and not bumpers["north"]:
        pos[0] -= 1
    elif action == "south" and not bumpers["south"]:
        pos[0] += 1
    elif action == "west" and not bumpers["west"]:
        pos[1] -= 1
    elif action == "east" and not bumpers["east"]:
        pos[1] += 1
    steps += 1
    print("Final grid:\n", grid)
    print("Total cleaned squares:", cleaned)

# Example grid
grid = np.array([
    [1, 0, 1],
    [0, 1, 0],
    [1, 0, 1]
])

agent = ModelBasedReflexAgent(grid.shape)
environment_model_based(agent, grid, max_steps=20)

```

Output:

Step 1: Pos [0, 0], Dirty: True, Action: suck

Step 2: Pos [0, 0], Dirty: False, Action: east

Step 3: Pos [0, 1], Dirty: False, Action: south

Step 4: Pos [1, 1], Dirty: True, Action: suck

Step 5: Pos [1, 1], Dirty: False, Action: south

Step 6: Pos [2, 1], Dirty: False, Action: east

Step 7: Pos [2, 2], Dirty: True, Action: suck

Step 8: Pos [2, 2], Dirty: False, Action: north

Step 9: Pos [1, 2], Dirty: False, Action: north

Step 10: Pos [0, 2], Dirty: True, Action: suck

Step 11: Pos [0, 2], Dirty: False, Action: south

Step 12: Pos [1, 2], Dirty: False, Action: north

Step 13: Pos [0, 2], Dirty: False, Action: west

Step 14: Pos [0, 1], Dirty: False, Action: east

Step 15: Pos [0, 2], Dirty: False, Action: south

Step 16: Pos [1, 2], Dirty: False, Action: south

Step 17: Pos [2, 2], Dirty: False, Action: west

Step 18: Pos [2, 1], Dirty: False, Action: west

Step 19: Pos [2, 0], Dirty: True, Action: suck

Final grid:

[[0 0 0]

[0 0 0]

[0 0 0]]

Total cleaned squares: 5

Post Lab Tasks

Task:

Implement a Simple Reflex Agent that monitors the temperature in a residential area using sensors. Implement a program in Python where the agent gets temperature readings from multiple sensors (assume 9 sensors are installed in different locations). Each sensor measures the temperature in Celsius.

- The agent should calculate the average temperature for all the sensors and convert it into Fahrenheit.
- Implement simple condition-action rules for deciding the actions based on the temperature range.

- o If the average temperature is above 85°F, the agent should activate a cooling mechanism
 - o if average temperature is below 60°F, the agent should activate heating.
 - o If the temperature is within the comfortable range (60°F to 85°F), no action should be taken.
- Visualize the sensor readings and the actions taken by the agent. Use a simple plot (e.g., using Matplotlib) to:
- o Display the temperature data for each sensor.
 - o Show the corresponding actions (cooling, heating, or no action).

Solution :

```

import numpy as np
import matplotlib.pyplot as plt

# Generate random Celsius readings for 9 sensors (e.g., between 10°C and 35°C)
sensor_readings_c = np.random.uniform(10, 35, 9)

# Calculate average in Celsius and convert to Fahrenheit
avg_c = np.mean(sensor_readings_c)
avg_f = avg_c * 9/5 + 32

# Decide action based on average Fahrenheit temperature
if avg_f > 85:
    action = "Cooling Activated"
elif avg_f < 60:
    action = "Heating Activated"
else:
    action = "No Action"

# Print results
print("Sensor readings (°C):", np.round(sensor_readings_c, 2))
print(f"Average temperature: {avg_c:.2f}°C / {avg_f:.2f}°F")
print("Action:", action)

# Visualization
plt.figure(figsize=(8,4))
plt.bar(range(1, 10), sensor_readings_c, color='skyblue')
plt.axhline((60-32)*5/9, color='orange', linestyle='--', label='60°F (Heating threshold)')
plt.axhline((85-32)*5/9, color='red', linestyle='--', label='85°F (Cooling threshold)')
plt.title(f"Sensor Temperatures and Agent Action: {action}")
plt.xlabel("Sensor Number")
plt.ylabel("Temperature (°C)")

```

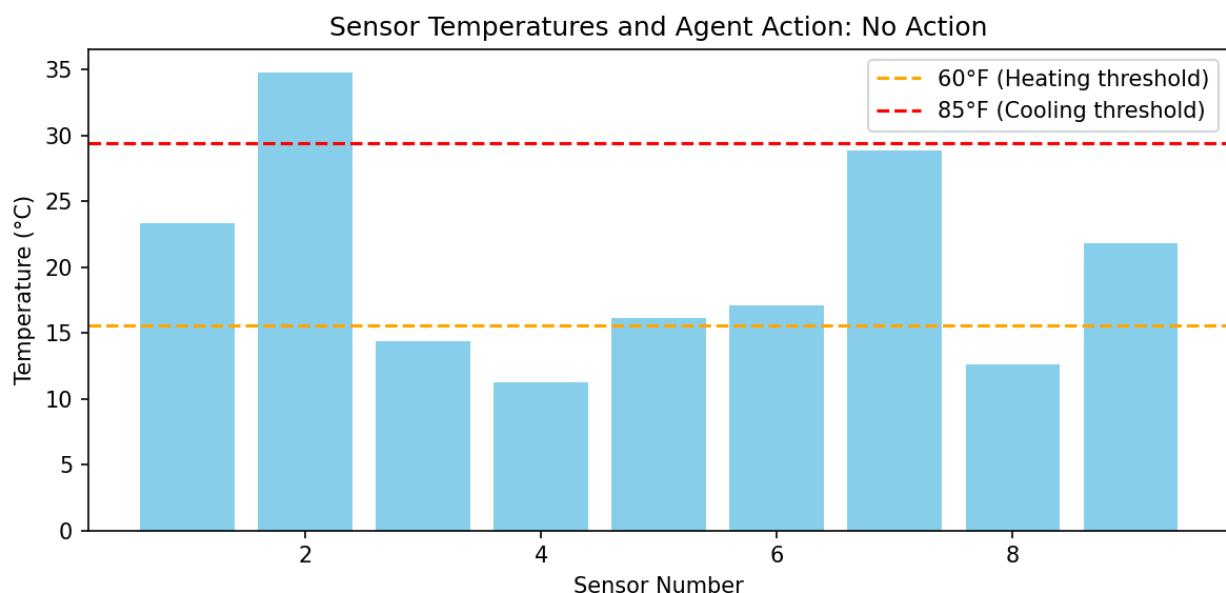
```
plt.legend()  
plt.tight_layout()  
plt.show()
```

Output:

Sensor readings (°C): [23.3 34.8 14.38 11.29 16.14 17.12 28.83 12.6 21.79]

Average temperature: 20.03°C / 68.05°F

Action: No Action



Task 4:

After successfully implementing a reflex agent, extend it to a Model-Based Agent. In your implementation, the agent should:

- ❖ Maintain an internal model of the environment (e.g., temperature history).
- ❖ Use this model to predict the future state of the environment and make
- ❖ decisions based on the predicted state, not just the current percept.

Solution :

```
import numpy as np  
import matplotlib.pyplot as plt  
  
class ModelBasedTempAgent:  
    def __init__(self):
```

```

    self.history = [] # Store temperature history

def update_history(self, readings_c):
    self.history.append(readings_c.copy())

def predict_next_avg(self):
    # Simple prediction: average of last 2 readings, or current if only one
    if len(self.history) >= 2:
        last = np.mean(self.history[-1])
        prev = np.mean(self.history[-2])
        return (last + prev) / 2
    elif self.history:
        return np.mean(self.history[-1])
    else:
        return None

def decide_action(self, predicted_avg_f):
    if predicted_avg_f > 85:
        return "Cooling Activated"
    elif predicted_avg_f < 60:
        return "Heating Activated"
    else:
        return "No Action"

# Simulate readings for 9 sensors over 10 time steps
agent = ModelBasedTempAgent()
actions = []
predicted_avgs_f = []

for t in range(10):
    readings_c = np.random.uniform(10, 35, 9)
    agent.update_history(readings_c)
    predicted_avg_c = agent.predict_next_avg()
    predicted_avg_f = predicted_avg_c * 9/5 + 32
    action = agent.decide_action(predicted_avg_f)
    actions.append(action)
    predicted_avgs_f.append(predicted_avg_f)
    print(f"Time {t+1}: Readings (°C): {np.round(readings_c,2)} | Predicted Avg: {predicted_avg_c:.2f}°C / {predicted_avg_f:.2f}°F | Action: {action}")

# Visualization
plt.figure(figsize=(10,5))
for i in range(9):
    plt.plot([agent.history[t][i] for t in range(10)], label=f"Sensor {i+1}")

```

```

plt.axhline((60-32)*5/9, color='orange', linestyle='--', label='60°F (Heating threshold)')
plt.axhline((85-32)*5/9, color='red', linestyle='--', label='85°F (Cooling threshold)')
plt.title("Sensor Temperatures Over Time")
plt.xlabel("Time Step")
plt.ylabel("Temperature (°C)")
plt.legend(loc='upper left', bbox_to_anchor=(1,1))
plt.tight_layout()
plt.show()

plt.figure(figsize=(8,3))
plt.plot(predicted_avgs_f, marker='o', label='Predicted Avg (°F)')
plt.title("Predicted Average Temperature (°F) and Actions")
plt.xlabel("Time Step")
plt.ylabel("Predicted Avg Temp (°F)")
for i, act in enumerate(actions):
    plt.text(i, predicted_avgs_f[i]+1, act, ha='center', fontsize=8)
plt.axhline(60, color='orange', linestyle='--')
plt.axhline(85, color='red', linestyle='--')
plt.tight_layout()
plt.show()

```

Output:

Time 1: Readings (°C): [27.07 29.9 30.31 21.82 32.6 26.21 14.75 16.06 24.14] | Predicted Avg: 24.76°C / 76.57°F | Action: No Action

Time 2: Readings (°C): [33.87 30.83 33.02 24.03 22.81 23.18 25.54 10.39 17.94] | Predicted Avg: 24.69°C / 76.45°F | Action: No Action

Time 3: Readings (°C): [15.36 29.8 16.76 29.58 23.23 32.26 30.43 30.36 33.49] | Predicted Avg: 25.72°C / 78.29°F | Action: No Action

Time 4: Readings (°C): [25. 12.9 16.95 29.02 24.47 26.88 29.29 29.37 17.35] | Predicted Avg: 25.14°C / 77.25°F | Action: No Action

Time 5: Readings (°C): [34.82 28.97 10.14 18.55 23.36 12.99 12.82 16.41 12.02] | Predicted Avg: 21.18°C / 70.13°F | Action: No Action

Time 6: Readings (°C): [26.21 16.91 22.85 17.27 12.24 29.53 25.87 21.79 12.79] | Predicted Avg: 19.75°C / 67.55°F | Action: No Action

Time 7: Readings (°C): [11.45 15.99 10.43 13.46 10.96 21.66 25.03 21.17 19.55] | Predicted Avg: 18.62°C / 65.52°F | Action: No Action

Time 8: Readings (°C): [33.22 27.41 20.96 20.73 19.04 15.46 21.59 22.66 16.08] | Predicted Avg: 19.27°C / 66.69°F | Action: No Action

Time 9: Readings (°C): [26.74 11.98 25.09 26.64 28.87 31.57 23.33 22.81 30.86] | Predicted Avg: 23.61°C / 74.50°F | Action: No Action

Time 10: Readings (°C): [27.63 29.23 20.09 22.89 30.65 15.97 19.64 24.01 32.14] | Predicted Avg: 25.01°C / 77.01°F | Action: No Action

