



## **NATIONAL UNIVERSITY OF SCIENCE AND TECHNOLOGY**

### **DEPARTMENT OF COMPUTER SCIENCE**

#### **ARTIFICIAL INTELLIGENCE LAB**

<b>NAME</b>	Ayesha Imran
<b>Class</b>	CS-A
<b>Oel</b>	08
<b>Course</b>	Artificial Intelligence
<b>Date</b>	9-December-25
<b>Submitted To</b>	Lec. Ijlal Haider

## IN LAB TASKS

**TASK 01:** For the given dataset “House Price Prediction Dataset.csv”, which contains house-related features influencing property prices, complete the following tasks:

- Data Exploration and Preprocessing:
- Perform exploratory data analysis (EDA) to:
- Understand the structure of the dataset (e.g., data types, missing values, and summary statistics).
- Visualize relationships between the features (Area, Bedrooms, Bathrooms, etc.) and the target variable (Price) using scatter plots, heatmaps, or box plots.
- Preprocess the data by:
- Handling missing values, if any.
- Encoding categorical variables like Condition and Location.
- Normalizing or scaling numerical features like Area and Year Built to ensure comparability across models.
- Regression Model Implementation:
- Develop models to predict house prices using the following regression techniques:
- Simple Linear Regression (e.g., predicting Price using Area alone).
- Multiple Linear Regression (using all numerical features).
- Ridge Regression
- Lasso Regression
- Decision Tree Regression
- Random Forest Regression
- Feature Engineering:
- Create new features to capture interactions or non-linear relationships, such as:
  - Interaction terms like Area x Condition or Bedrooms x Bathrooms.
  - Polynomial features (e.g.,  $\text{Area}^2$ ) for non-linear regression.
- Performance Evaluation:
- Evaluate the performance of each model using:
- Mean Absolute Error (MAE)

- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- Coefficient of Determination (R<sup>2</sup>)
- Comparison and Insights:
- Compare the models' performance metrics to determine:
- Which features have the strongest influence on Price predictions.
- Which model provides the most accurate and reliable results.
- Use visualization tools to show Actual vs. Predicted Price for each model.

**CODE:**

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.feature_selection import SelectKBest, f_regression
import warnings
warnings.filterwarnings('ignore')

# =====
# 1. Load and Explore the Data
# =====
df = pd.read_csv('House Prediction Dataset.csv') # Change path if needed

print("Dataset Shape:", df.shape)
print("\nData Types:\n", df.dtypes)
print("\nMissing Values:\n", df.isnull().sum())
print("\nSummary Statistics:\n", df.describe())

# Target and features
target = 'price'
numerical_features = ['area', 'bedrooms', 'bathrooms', 'stories', 'parking']
categorical_features = ['mainroad', 'guestroom', 'basement', 'hotwaterheating',
                       'airconditioning', 'prefarea', 'furnishingstatus']

```

```

# =====
# 2. Visualizations
# =====
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()

# Scatter plots
plt.figure(figsize=(12, 6))
sns.scatterplot(data=df, x='area', y='price', hue='furnishingstatus')
plt.title('Price vs Area colored by Furnishing Status')
plt.show()

# Box plots
plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x='furnishingstatus', y='price')
plt.title('Price Distribution by Furnishing Status')
plt.show()

# =====
# 3. Feature Engineering
# =====
# Interaction terms
df['area_bedrooms'] = df['area'] * df['bedrooms']
df['bedrooms_bathrooms'] = df['bedrooms'] * df['bathrooms']
df['area_furnished'] = df['area'] * (df['furnishingstatus'] ==
'furnished').astype(int)

# Polynomial feature (area squared)
df['area_squared'] = df['area'] ** 2

# Update numerical features list
numerical_features += ['area_bedrooms', 'bedrooms_bathrooms', 'area_furnished',
'area_squared']

# =====
# 4. Data Preprocessing Pipeline
# =====
X = df.drop(target, axis=1)
y = df[target]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

```

```
# Preprocessing
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(drop='first', handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_features),
        ('cat', categorical_transformer, categorical_features)
    ]
)

# =====
# 5. Models
# =====

models = {
    'Simple Linear (Area only)': Pipeline([
        ('preprocessor', ColumnTransformer([('num', StandardScaler(), ['area'])])),
        ('regressor', LinearRegression())
]),

    'Multiple Linear': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', LinearRegression())
]),

    'Ridge': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', Ridge(alpha=1.0))
]),

    'Lasso': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', Lasso(alpha=1.0))
]),

    'Decision Tree': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', DecisionTreeRegressor(max_depth=5, random_state=42))
])}
```

```

'Random Forest': Pipeline([
    ('preprocessor', preprocessor),
    ('regressor', RandomForestRegressor(n_estimators=200, max_depth=10,
random_state=42))
])
}

# =====
# 6. Train and Evaluate
# =====
results = []

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    results.append({
        'Model': name,
        'MAE': round(mae, 0),
        'MSE': round(mse, 0),
        'RMSE': round(rmse, 0),
        'R2

```

```

print("\nModel Performance Comparison:")
print(results_df.sort_values('R²', ascending=False))

# =====
# 8. Feature Importance (for Random Forest)
# =====
rf_model = models['Random Forest']
rf_model.fit(X_train, y_train)

# Get feature names after transformation
feature_names = (rf_model.named_steps['preprocessor']
                  .get_feature_names_out())

importances = rf_model.named_steps['regressor'].feature_importances_
feat_imp = pd.Series(importances,
                     index=feature_names).sort_values(ascending=False)

plt.figure(figsize=(12, 8))
sns.barplot(x=feat_imp.values, y=feat_imp.index)
plt.title('Feature Importance (Random Forest)')
plt.show()

```

## Output:

```

Dataset Shape: (2000, 10)

Data Types:
   Id          int64
   Area        int64
   Bedrooms    int64
   Bathrooms   int64
   Floors       int64
   YearBuilt    int64
   Location     object
   Condition    object
   Garage       object
   Price         int64
dtype: object

```

```
Missing Values:
```

```
    Id      0
    Area     0
    Bedrooms 0
    Bathrooms 0
    Floors    0
    YearBuilt 0
    Location   0
    Condition   0
    Garage     0
    Price      0
dtype: int64
```


```
Summary Statistics:
```

```
          Id      Area    Bedrooms ...    Floors    YearBuilt    Price
count  2000.000000 2000.000000 2000.000000 ... 2000.000000 2000.000000 2000.000000
mean   1000.500000 2786.209500    3.003500 ... 1.993500 1961.446000 537676.855000
std    577.494589 1295.146799    1.424606 ... 0.809188 35.926695 276428.845719
Price      0
dtype: int64
```

```
Summary Statistics:
```

```
          Id      Area    Bedrooms ...    Floors    YearBuilt    Price
count  2000.000000 2000.000000 2000.000000 ... 2000.000000 2000.000000 2000.000000
mean   1000.500000 2786.209500    3.003500 ... 1.993500 1961.446000 537676.855000
std    577.494589 1295.146799    1.424606 ... 0.809188 35.926695 276428.845719
          Id      Area    Bedrooms ...    Floors    YearBuilt    Price
count  2000.000000 2000.000000 2000.000000 ... 2000.000000 2000.000000 2000.000000
mean   1000.500000 2786.209500    3.003500 ... 1.993500 1961.446000 537676.855000
std    577.494589 1295.146799    1.424606 ... 0.809188 35.926695 276428.845719
          mean  1000.500000 2786.209500    3.003500 ... 1.993500 1961.446000 537676.855000
          std   577.494589 1295.146799    1.424606 ... 0.809188 35.926695 276428.845719
          min   1.000000 501.000000    1.000000 ... 1.000000 1900.000000 50005.000000
          25%  500.750000 1653.000000    2.000000 ... 1.000000 1930.000000 300098.000000
          25%  500.750000 1653.000000    2.000000 ... 1.000000 1930.000000 300098.000000
          50% 1000.500000 2833.000000    3.000000 ... 2.000000 1961.000000 539254.000000
          75% 1500.250000 3887.500000    4.000000 ... 3.000000 1993.000000 780086.000000
          max 2000.000000 4999.000000    5.000000 ... 3.000000 2023.000000 999656.000000
```

```
[8 rows x 7 columns]
```



## POST LAB TASKS

**TASK 02:** For the given dataset For the give dataset “Advertising.csv”, which contains information about advertising spend and its impact on sales:

Data Exploration and Preprocessing:

- Perform exploratory data analysis (EDA) to understand the structure of the dataset, including statistical summaries and visualizations to explore relationships between features (TV, Radio, Newspaper) and the target variable (Sales).
- Handle any missing values if present and apply necessary preprocessing steps, such as feature scaling or normalization, to prepare the data for modeling.

Regression Model Implementation:

- Apply the following regression models to predict Sales based on advertising

spending:

- Simple Linear Regression (e.g., using one predictor like TV or Radio).
- Multiple Linear Regression (using all predictors: TV, Radio, and Newspaper).
- Ridge Regression
- Lasso Regression
- Decision Tree Regression
- Random Forest Regression

Performance Evaluation:

1. Evaluate each models performance using the following metrics:
2. Mean Absolute Error (MAE)
3. Mean Squared Error (MSE)
4. Root Mean Squared Error (RMSE)
5. Coefficient of Determination (R<sup>2</sup>)
6. Comparison and Insights:
7. Compare the results of all models based on the evaluation metrics.
8. Summarize your findings, including:
  - Which advertising medium (TV, Radio, or Newspaper) has the most significant impact on Sales.
  - Which regression model performs best for this dataset and why.
  - Any potential overfitting or underfitting observed during the analysis.

## CODE:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import warnings
warnings.filterwarnings('ignore')
```

```

# =====
# 1. Load and Explore the Data
# =====
df = pd.read_csv('Advertising.csv') # Change path if needed

print("Dataset Shape:", df.shape)
print("\nFirst 5 rows:\n", df.head())
print("\nMissing Values:\n", df.isnull().sum())
print("\nSummary Statistics:\n", df.describe())

# Features and target
features = ['TV', 'Radio', 'Newspaper']
target = 'Sales'

# =====
# 2. Visualizations
# =====
# Pairplot to see relationships
sns.pairplot(df, x_vars=features, y_vars=target, height=4, aspect=1,
kind='scatter')
plt.suptitle('Advertising Spend vs Sales', y=1.02)
plt.show()

# Correlation heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', fmt='.3f')
plt.title('Correlation Matrix')
plt.show()

# Boxplot for outliers
plt.figure(figsize=(10, 6))
sns.boxplot(data=df[features])
plt.title('Distribution of Advertising Spend')
plt.show()

# =====
# 3. Preprocessing
# =====
# No missing values in this dataset (confirmed)

X = df[features]
y = df[target]

# Train-test split (80-20)

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Scale features (good practice for Ridge/Lasso and some tree models)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# =====
# 4. Models
# =====

models = {
    'Simple Linear (TV only)': LinearRegression(),
    'Multiple Linear': LinearRegression(),
    'Ridge': Ridge(alpha=1.0),
    'Lasso': Lasso(alpha=1.0),
    'Decision Tree': DecisionTreeRegressor(max_depth=5, random_state=42),
    'Random Forest': RandomForestRegressor(n_estimators=200, max_depth=10,
random_state=42)
}

# For Simple Linear (TV only), we use only TV
X_train_tv = X_train[['TV']]
X_test_tv = X_test[['TV']]
X_train_tv_scaled = scaler.fit_transform(X_train_tv)
X_test_tv_scaled = scaler.transform(X_test_tv)

results = []

for name, model in models.items():
    if name == 'Simple Linear (TV only)':
        model.fit(X_train_tv_scaled, y_train)
        y_pred = model.predict(X_test_tv_scaled)
    else:
        model.fit(X_train_scaled, y_train)
        y_pred = model.predict(X_test_scaled)

    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    results.append({
        'Model': name,
        'MAE': round(mae, 2),
        'MSE': round(mse, 2),
        'RMSE': round(rmse, 2),
        'R2': round(r2, 2)
    })
}

```

```

        'MSE': round(mse, 2),
        'RMSE': round(rmse, 2),
        'R22', ascending=False))

# =====
# 6. Feature Importance (for Random Forest)
# =====
rf_model = models['Random Forest']
rf_model.fit(X_train_scaled, y_train)

feat_imp = pd.Series(rf_model.feature_importances_,
index=features).sort_values(ascending=False)

plt.figure(figsize=(8, 5))
sns.barplot(x=feat_imp.values, y=feat_imp.index, palette='viridis')
plt.title('Feature Importance - Random Forest')
plt.xlabel('Importance Score')
plt.show()

# =====
# 7. Coefficients for Linear Models
# =====
print("\nCoefficients from Multiple Linear Regression:")
lr = LinearRegression()
lr.fit(X_train_scaled, y_train)
coef = pd.Series(lr.coef_, index=features).sort_values(ascending=False)

```



```
print(coef)
```

## OUTPUT:

```
First 5 rows:  
      Unnamed: 0      TV  Radio  Newspaper  Sales  
0            1  230.1   37.8      69.2   22.1  
1            2   44.5   39.3      45.1   10.4  
2            3   17.2   45.9      69.3    9.3  
3            4  151.5   41.3      58.5   18.5  
4            5  180.8   10.8      58.4   12.9  
  
Missing Values:  
      Unnamed: 0      0  
      TV          0  
      Radio        0  
      Newspaper    0  
      Sales        0  
      dtype: int64  
  
Summary Statistics:  
Missing Values:  
      Unnamed: 0      0  
      TV          0  
      Radio        0  
      Newspaper    0  
      Sales        0  
      dtype: int64
```

```
Summary Statistics:  
TV      0  
Radio   0  
Newspaper 0  
Sales    0  
dtype: int64  
  
Summary Statistics:  
Radio   0  
Newspaper 0  
Sales    0  
dtype: int64  
  
Summary Statistics:  
Newspaper 0  
Sales    0  
dtype: int64  
  
Summary Statistics:  
dtype: int64  
  
Summary Statistics:  
  
Summary Statistics:  
    Unnamed: 0          TV        Radio    Newspaper       Sales  
    Unnamed: 0          TV        Radio    Newspaper       Sales  
count  200.000000  200.000000  200.000000  200.000000  200.000000
```

```
Summary Statistics:  
    Unnamed: 0          TV        Radio    Newspaper       Sales  
    Unnamed: 0          TV        Radio    Newspaper       Sales  
count  200.000000  200.000000  200.000000  200.000000  200.000000  
mean   100.500000  147.042500  23.264000  30.554000  14.022500  
std    57.879185  85.854236  14.846809  21.778621  5.217457  
mean   100.500000  147.042500  23.264000  30.554000  14.022500  
std    57.879185  85.854236  14.846809  21.778621  5.217457  
min    1.000000   0.700000   0.000000   0.300000   1.600000  
std    57.879185  85.854236  14.846809  21.778621  5.217457  
min    1.000000   0.700000   0.000000   0.300000   1.600000  
min    1.000000   0.700000   0.000000   0.300000   1.600000  
25%    50.750000  74.375000  9.975000  12.750000  10.375000  
50%   100.500000  149.750000  22.900000  25.750000  12.900000  
75%   150.250000  218.825000  36.525000  45.100000  17.400000  
max   200.000000  296.400000  49.600000  114.000000  27.000000
```



## Task 2

For the give dataset “Store\_CA.csv”, which contains key performance indicators (KPIs) related to retail store operations, customer engagement, and sales performance:

Exploratory Data Analysis (EDA) and Preprocessing:

- Analyze the dataset to understand the structure and relationships between features, including visualizations and statistical summaries.
- Handle missing values, if any, and preprocess data for modeling.
- Normalize or scale numerical features
- StoreSize, MarketingSpend, CustomerFootfall, etc.
- Encode categorical features such as,
- StoreLocation and StoreCategory appropriately.
- Explore correlations between features and the target variable, MonthlySalesRevenue.
- Regression Model Implementation:
- Build and evaluate the following regression models:
- Simple Linear Regression
- Like MarketingSpend vs. MonthlySalesRevenue
- Multiple Linear Regression (consider multiple predictors)
- Like MarketingSpend, CustomerFootfall, etc.)
- Ridge Regression
- Lasso Regression
- Decision Tree Regression
- Random Forest Regression
- Model Evaluation:
- Evaluate each model using the following metrics:
- Mean Absolute Error (MAE)
- Mean Squared Error (MSE)
- Root Mean Squared Error (RMSE)
- Coefficient of Determination (R<sup>2</sup>)
- Comparison and Insights:
- Compare the models based on their performance metrics.
- Discuss findings, highlighting:
- Which factors have the strongest influence on MonthlySalesRevenue.
- Which regression model is most suitable for this dataset and why.

## CODE:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import warnings
warnings.filterwarnings('ignore')

# =====
# 1. Load and Explore the Data
# =====
df = pd.read_csv('Store CA.csv') # Change path if needed

print("Dataset Shape:", df.shape)
print("\nColumns:", df.columns.tolist())
print("\nFirst 5 rows:")
print(df.head())
print("\nMissing Values:")
print(df.isnull().sum())
print("\nSummary Statistics:")
print(df.describe().round(2))

# Identify target and features
target = 'MonthlySalesRevenue'

# Automatically detect numerical and categorical columns
numerical_cols = df.select_dtypes(include=['int64',
'float64']).columns.drop(target, errors='ignore').tolist()
categorical_cols = df.select_dtypes(include=['object',
'category']).columns.tolist()

print("\nNumerical Features:", numerical_cols)
print("Categorical Features:", categorical_cols)

# =====
# 2. Visualizations & EDA
# =====
# Correlation heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm', fmt='.2f')
plt.title('Correlation Heatmap')
plt.show()
```

```

# Pairplot for key features vs target
key_features = [col for col in numerical_cols if col in ['MarketingSpend',
'CustomerFootfall', 'StoreSize', 'EmployeeCount']]
sns.pairplot(df, x_vars=key_features, y_vars=target, height=4, aspect=1)
plt.suptitle('Key Features vs Monthly Sales Revenue', y=1.02)
plt.show()

# Boxplot for categorical features
if 'StoreLocation' in df.columns:
    plt.figure(figsize=(10, 6))
    sns.boxplot(data=df, x='StoreLocation', y=target)
    plt.title('Monthly Sales Revenue by Store Location')
    plt.xticks(rotation=45)
    plt.show()

if 'StoreCategory' in df.columns:
    plt.figure(figsize=(10, 6))
    sns.boxplot(data=df, x='StoreCategory', y=target)
    plt.title('Monthly Sales Revenue by Store Category')
    plt.xticks(rotation=45)
    plt.show()

# =====
# 3. Preprocessing
# =====
X = df.drop(target, axis=1)
y = df[target]

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Preprocessing pipeline
numeric_transformer = Pipeline(steps=[
    ('scaler', StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ('onehot', OneHotEncoder(drop='first', handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numerical_cols),
        ('cat', categorical_transformer, categorical_cols)
])

```

```
])

# =====
# 4. Models
# =====

models = {
    'Simple Linear (MarketingSpend only)': Pipeline([
        ('preprocessor', ColumnTransformer([('num', StandardScaler(), [
            'MarketingSpend'])])),
        ('regressor', LinearRegression())
    ]),

    'Multiple Linear Regression': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', LinearRegression())
    ]),

    'Ridge Regression': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', Ridge(alpha=1.0))
    ]),

    'Lasso Regression': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', Lasso(alpha=1.0))
    ]),

    'Decision Tree': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', DecisionTreeRegressor(max_depth=8, random_state=42))
    ]),

    'Random Forest': Pipeline([
        ('preprocessor', preprocessor),
        ('regressor', RandomForestRegressor(n_estimators=200, max_depth=10,
random_state=42))
    ])
}

results = []

for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
```

```

mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

results.append({
    'Model': name,
    'MAE': round(mae, 0),
    'MSE': round(mse, 0),
    'RMSE': round(rmse, 0),
    'R22', ascending=False))

# =====
# 6. Feature Importance (Random Forest)
# =====
rf_model = models['Random Forest']
rf_model.fit(X_train, y_train)

# Get feature names after transformation
feature_names = rf_model.named_steps['preprocessor'].get_feature_names_out()

importances = rf_model.named_steps['regressor'].feature_importances_
feat_imp = pd.Series(importances,
index=feature_names).sort_values(ascending=False)

plt.figure(figsize=(12, 8))

```

```

sns.barplot(x=feat_imp.values, y=feat_imp.index)
plt.title('Feature Importance - Random Forest')
plt.xlabel('Importance Score')
plt.show()

# =====
# 7. Insights Summary
# =====
print("\n==== Key Insights ===")
print("Strongest predictors of MonthlySalesRevenue:")
print(feat_imp.head(8)) # Top 8 features

best_model = results_df.loc[results_df['R²'].idxmax()]
print(f"\nBest performing model: {best_model['Model']} ")
print(f"R² = {best_model['R²']}, RMSE = {best_model['RMSE']} ")
print("\nWhy Random Forest usually wins:")
print("- Handles non-linear relationships and interactions well")
print("- Robust to outliers")
print("- Provides feature importance directly")
print("- Less prone to overfitting than Decision Tree")

```

## OUTPUT:

```

Dataset Shape: (1650, 12)

Columns: ['ProductVariety', 'MarketingSpend', 'CustomerFootfall', 'StoreSize', 'EmployeeEfficiency', 'StoreAge', 'CompetitorDistance', 'PromotionsCount', 'EconomicIndicator', 'StoreLocation', 'StoreCategory', 'MonthlySalesRevenue']

First 5 rows:
   ProductVariety  MarketingSpend  ...  StoreCategory  MonthlySalesRevenue
0            581              29  ...    Electronics      284.90
1            382              31  ...    Electronics      308.21
2            449              35  ...       Grocery      292.11
3            666               9  ...     Clothing      279.61
4            657              35  ...    Electronics      359.71

[5 rows x 12 columns]


Missing Values:
ProductVariety      0
MarketingSpend      0
CustomerFootfall    0
StoreSize            0
EmployeeEfficiency  0
StoreAge             0
CompetitorDistance  0
PromotionsCount     0
EconomicIndicator   0
StoreLocation        0
StoreCategory        0

```

Ln 4, Col 1 (6678 selected) Spaces: 4 UTF-8 CRLF

Summary Statistics:										
	ProductVariety	MarketingSpend	...	EconomicIndicator	MonthlySalesRevenue					
count	1650.00	1650.00	...	1650.00	1650.00					
mean	500.92	27.46	...	99.76	299.25					
std	148.05	13.01	...	14.61	65.54					
min	100.00	5.00	...	60.00	106.71					
min	100.00	5.00	...	60.00	106.71					
min	100.00	5.00	...	60.00	106.71					
min	100.00	5.00	...	60.00	106.71	5%	396.00			
17.00	...	90.00		254.12						
25%	396.00	17.00	...	90.00	254.12	5%	602.75			
38.00	...	109.60		344.22						
50%	500.50	27.00	...	100.30	297.44	min	100.00			
5.00	...	60.00		106.71						
25%	396.00	17.00	...	90.00	254.12					
min	100.00	5.00	...	60.00	106.71					
25%	396.00	17.00	...	90.00	254.12					
50%	500.50	27.00	...	100.30	297.44	min	100.00			
5.00	...	60.00		106.71						
25%	396.00	17.00	...	90.00	254.12					
min	100.00	5.00	...	60.00	106.71					
25%	396.00	17.00	...	90.00	254.12					
50%	500.50	27.00	...	100.30	297.44					
75%	602.75	38.00	...	109.60	344.22					
max	1092.00	50.00	...	140.00	534.26					
[ 8 rows x 10 columns ]										

Numerical Features: ['ProductVariety', 'MarketingSpend', 'CustomerFootfall', 'StoreSize', 'EmployeeEfficiency', 'StoreAge', 'CompetitorDistance', 'PromotionsCount', 'EconomicIndicator']  
Categorical Features: ['StoreLocation', 'StoreCategory']



*END*