



## **NATIONAL UNIVERSITY OF SCIENCE AND TECHNOLOGY**

### **DEPARTMENT OF COMPUTER SCIENCE**

#### **ARTIFICIAL INTELLIGENCE LAB**

<b>NAME</b>	Ayesha Imran
<b>Class</b>	CS-A
<b>Oel</b>	OEL 02
<b>Course</b>	Artificial Intelligence
<b>Date</b>	29 <sup>th</sup> -December-25
<b>Submitted To</b>	Lec. Ijlal Haider

# IN LAB TASKS

## TASK 01: Linear Regression Model:

```
# salary_Linear_regression.py
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# -----
# LOAD DATA
# -----
df = pd.read_csv('Salary Data.csv')

print("Dataset shape:", df.shape)
print(df.columns.tolist())

# -----
# PREPROCESSING
# -----
df = df.dropna().reset_index(drop=True)

# Encode categorical columns
cat_columns = ['Gender', 'Education Level', 'Job Title']
label_enc = {}
for col in cat_columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_enc[col] = le # save if you want to decode later

# Features & target
X = df.drop('Salary', axis=1)
y = df['Salary']

# Scale features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train / Test split
X_train, X_test, y_train, y_test = train_test_split(
```

```

        X_scaled, y, test_size=0.2, random_state=42
    )

print(f"Train: {X_train.shape[0]} samples | Test: {X_test.shape[0]} samples")

# -----
# MODEL: Linear Regression
# -----
model = LinearRegression()
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

# -----
# EVALUATION
# -----
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("\nLINEAR REGRESSION RESULTS")
print("-----")
print(f"Mean Squared Error : {mse:.0f}")
print(f"R2 Score : {r2:.4f}")
print(f"Adjusted R2 (approx): {1 - (1 - r2) * (len(y_test) - 1) / (len(y_test) - X_test.shape[1] - 1):.4f}")

# -----
# VISUALIZATION
# -----
plt.figure(figsize=(9, 6))
plt.scatter(y_test, y_pred, alpha=0.6, color='blue', label='Predictions')
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'r--', lw=2, label='Perfect fit')
plt.xlabel('Actual Salary')
plt.ylabel('Predicted Salary')
plt.title('Linear Regression - Actual vs Predicted Salary')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

## OUTPUT:

```
Dataset shape: (6704, 6)
['Age', 'Gender', 'Education Level', 'Job Title', 'Years of Experience', 'Salary']
Train: 5358 samples | Test: 1340 samples

LINEAR REGRESSION RESULTS


Mean Squared Error : 926,443,019
Train: 5358 samples | Test: 1340 samples

LINEAR REGRESSION RESULTS

LINEAR REGRESSION RESULTS

Mean Squared Error : 926,443,019
R2 Score : 0.6751
Adjusted R2 (approx): 0.6739
```

In [40]: Cell 27 - Session 4 - LITE - 8



## SVM Model :

```
# salary_svr.py
import pandas as pd
import numpy as np
```

```

import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, r2_score

# -----
# LOAD DATA
# -----
df = pd.read_csv('Salary Data.csv') # ← change path/name if needed

print("Dataset shape:", df.shape)
print(df.columns.tolist())

# -----
# PREPROCESSING
# -----
df = df.dropna().reset_index(drop=True)

# Encode categorical columns
cat_columns = ['Gender', 'Education Level', 'Job Title']
for col in cat_columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])

# Features & target
X = df.drop('Salary', axis=1)
y = df['Salary']

# IMPORTANT: SVR performs MUCH better with scaled data
scaler_X = StandardScaler()
X_scaled = scaler_X.fit_transform(X)

# Scale target too (helps a lot with SVR)
scaler_y = StandardScaler()
y_scaled = scaler_y.fit_transform(y.values.reshape(-1, 1)).ravel()

# Train / Test split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y_scaled, test_size=0.2, random_state=42
)

print(f"Train: {X_train.shape[0]} samples | Test: {X_test.shape[0]} samples")
# -----

```

```

# MODEL: Support Vector Regression (RBF kernel)
# -----
model = SVR(kernel='rbf', C=100, epsilon=0.1, gamma='scale')
model.fit(X_train, y_train)

# Predict (scaled)
y_pred_scaled = model.predict(X_test)

# Transform back to original salary scale
y_test_original = scaler_y.inverse_transform(y_test.reshape(-1, 1)).ravel()
y_pred_original = scaler_y.inverse_transform(y_pred_scaled.reshape(-1,
1)).ravel()

# -----
# EVALUATION
# -----
mse = mean_squared_error(y_test_original, y_pred_original)
r2 = r2_score(y_test_original, y_pred_original)

print("\nSUPPORT VECTOR REGRESSION (RBF) RESULTS")
print("-----")
print(f"Mean Squared Error : {mse:.4f}")
print(f"R2 Score : {r2:.4f}")

# -----
# VISUALIZATION
# -----
plt.figure(figsize=(9, 6))
plt.scatter(y_test_original, y_pred_original, alpha=0.6, color='darkred',
label='SVR Predictions')
plt.plot([y.min(), y.max()], [y.min(), y.max()], 'k--', lw=2, label='Perfect
fit')
plt.xlabel('Actual Salary')
plt.ylabel('Predicted Salary')
plt.title('SVR (RBF) - Actual vs Predicted Salary')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

## OUTPUT:

```

Dataset shape: (6704, 6)
['Age', 'Gender', 'Education Level', 'Job Title', 'Years of Experience', 'Salary']
Train: 5358 samples | Test: 1340 samples


```

#### SUPPORT VECTOR REGRESSION (RBF) RESULTS

```

Mean Squared Error : 235,441,521
R2 Score : 0.9174

```



#### Comparison:

CRITERION	LINEAR REGRESSION	SUPPORT VECTOR REGRESSION (SVR WITH RBF KERNEL)	WINNER (ON THIS DATASET)
MODEL TYPE	Linear	Non-linear (kernel trick)	—
TYPICAL R <sup>2</sup> SCORE	0.85 – 0.92	0.88 – 0.95 (often slightly higher)	SVR (usually)
TYPICAL MSE	Higher (e.g., 150M–300M)	Lower (e.g., 100M–220M)	SVR

<b>ABILITY TO CAPTURE NON-LINEARITY</b>	Poor – assumes straight-line relationships	Excellent – RBF kernel captures complex patterns	SVR
<b>HANDLING CATEGORICAL FEATURES AFTER ENCODING</b>	Good	Very good (especially with scaled data)	SVR (slight edge)
<b>SENSITIVITY TO FEATURE SCALING</b>	Moderate	Very high – must scale both X and y	—
<b>TRAINING SPEED</b>	Very fast	Slower (especially on >5k rows)	Linear Regression
<b>PREDICTION SPEED</b>	Extremely fast	Slower (depends on # of support vectors)	Linear Regression
<b>INTERPRETABILITY</b>	Excellent – you get coefficients for each feature	Poor – black-box model	Linear Regression
<b>OVERRFITTING RISK</b>	Low (especially with many features)	Higher if C is too large or gamma is too small	Linear Regression (safer)
<b>HYPERPARAMETERS TO TUNE</b>	Almost none	C, epsilon, gamma, kernel → needs tuning	Linear Regression (easier)
<b>BEST WHEN DATASET IS</b>	Medium-sized, mostly linear patterns	Medium/large, complex/non-linear patterns	—
<b>ROBUST TO OUTLIERS</b>	Sensitive	More robust (especially with epsilon-tube)	SVR

*END*

