

# Artificial Intelligence (CS280)

Noushin Saba

# **Lecture 3**

## **Reasoning and Knowledge Representation: A Theoretical Analysis**

# Outline

## Introduction to Knowledge and Reasoning

### Formal Foundations of Knowledge Representation

- Facts, Models, Entailment
- Knowledge based Agents

### Knowledge Representation Techniques

- Logic (Prepositional & First Order)
- Semantic Networks, Frames, Rules, Ontologies

### Reasoning Approaches

- Deductive, Inductive, Abductive
- Inference Procedures
- Soundness & Completeness

### Modern Views of KR

- Large Language Models (LLMs)
- Probabilistic Reasoning

# Outline

## Introduction to Knowledge and Reasoning

### Formal Foundations of Knowledge Representation

- Facts, Models, Entailment
- Knowledge based Agents

### Knowledge Representation Techniques

- Logic (Prepositional & First Order)
- Semantic Networks, Frames, Rules, Ontologies

### Reasoning Approaches

- Deductive, Inductive, Abductive
- Inference Procedures
- Soundness & Completeness

### Modern Views of KR

- Large Language Models (LLMs)
- Probabilistic Reasoning

# Introduction to Knowledge and Reasoning

- **What is Knowledge?**

- Knowledge = **facts + relationships + rules** about the world.
- Structured, contextual, and meaningful information.
  - Fact: “It is raining.”
  - Knowledge: “If it rains, the ground becomes wet.”

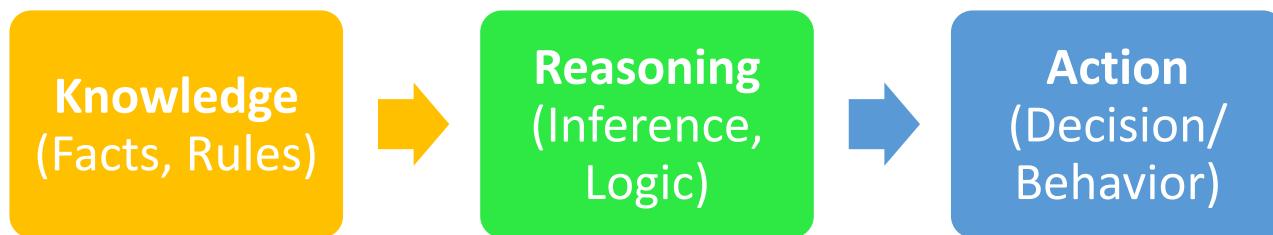
- **What is Reasoning?**

- The process of **using knowledge to draw conclusions** or make decisions.
  - Knowledge: “If it rains, the ground becomes wet.”
  - Fact: “It is raining.”
  - Reasoning → “The ground is wet.”

# Introduction to Knowledge and Reasoning (cont..)

- **Why Important in AI?**

- Forms the **foundation of intelligent agents**.
- Separates **data (knowledge)** from **processing (reasoning)**.
- Enables decision-making, problem-solving, and planning.



# Outline

Introduction  
to  
Knowledge  
and  
Reasoning

**Formal Foundations of Knowledge Representation**

- Facts, Models, Entailment
- Knowledge based Agents

**Knowledge Representation Techniques**

- Logic (Prepositional & First Order)
- Semantic Networks, Frames, Rules, Ontologies

**Reasoning Approaches**

- Deductive, Inductive, Abductive
- Inference Procedures
- Soundness & Completeness

**Modern Views of KR**

- Large Language Models (LLMs)
- Probabilistic Reasoning

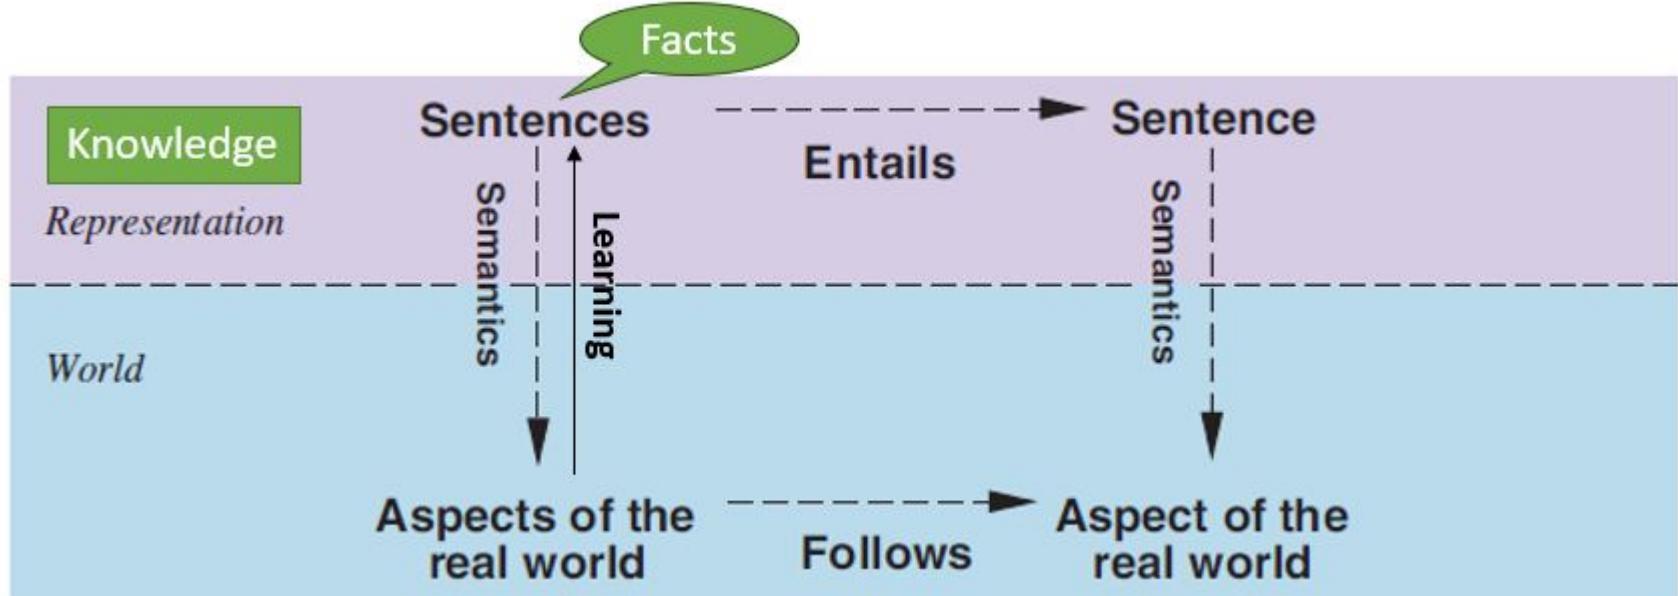
# Formal Foundations of Knowledge Representation

- **Facts**
  - Sentences that are known to be true.
- **Models (Possible Worlds)**
  - A **model** = a complete assignment of truth values to all facts.
  - **Possible worlds** = all models consistent with what we know.
  - Learning new facts reduces the number of possible worlds.
    - **Fact:** “It is raining.”
    - **Model 1:** Raining = True, Ground wet = True.
    - **Model 2:** Raining = False, Ground wet = False.
    - *Only Model 1 is consistent with the fact.*

# Formal Foundations of Knowledge Representation (cont..)

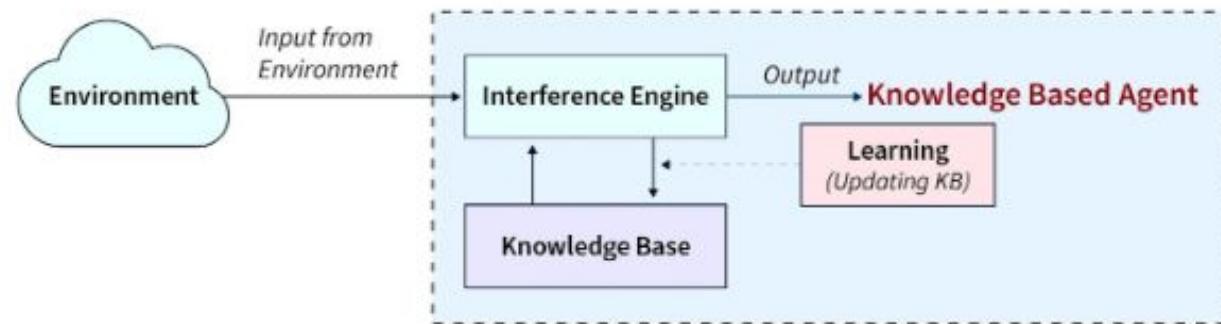
- **Entailment ( $\models$ )**

- A sentence **logically follows** from what we already know.
- If  $\text{KB} \models \alpha$ , then in every model where KB is true,  $\alpha$  is also true.
  - **KB** = “It is raining  $\rightarrow$  The ground is wet.”
  - **Fact** = “It is raining.”
  - **Entailment**: “The ground is wet.”



# Knowledge-Based Agents

- An **agent** that **reasons** and **acts** using a **Knowledge Base (KB)**.
- KB = **Set of facts** expressed in a formal language.
- **Inference Engine**: derives **new facts** from KB using **entailment**.
- **Advantages of KB Agents**
  - **Separation of Knowledge & Reasoning**: flexible design.
  - Can **learn new facts** and adapt actions.
  - Foundation of **expert systems**.



# Knowledge-Based Agents (cont..)

- **Key Components**

- **Knowledge Base (KB):**

- Stores facts, rules, and background knowledge.
    - Fact: “It is raining.”
    - Rule: “If it rains, the ground is wet.”

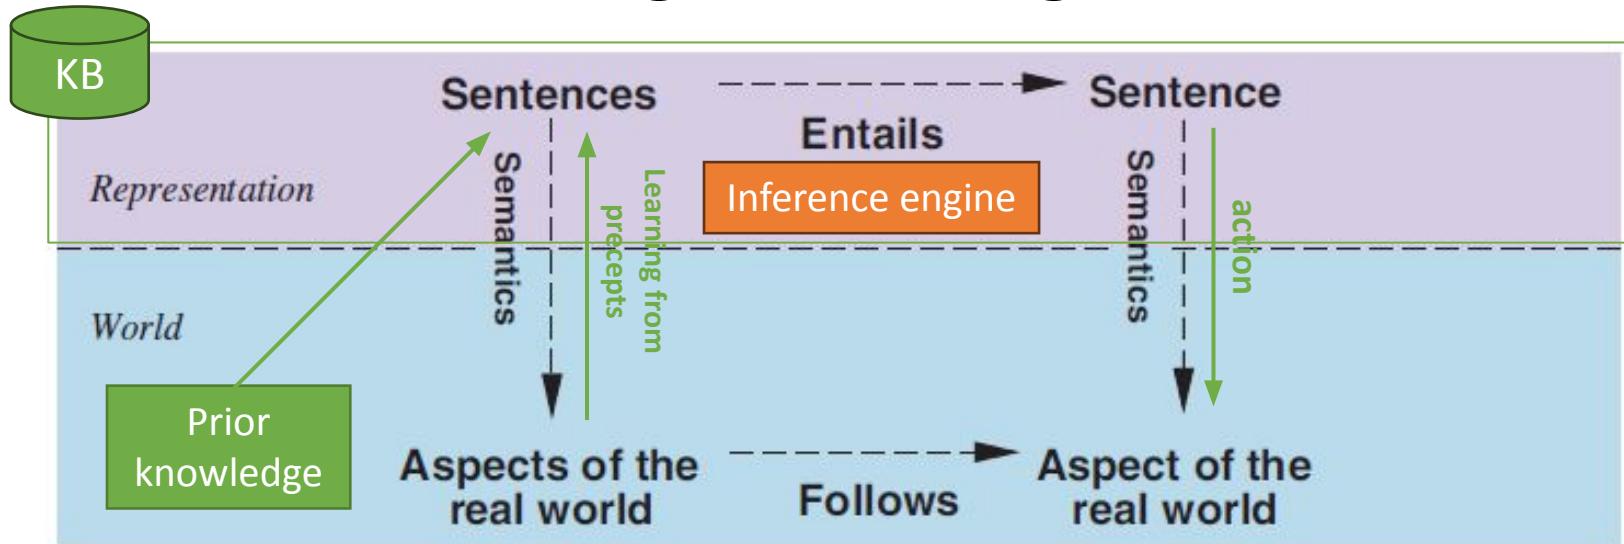
- **Inference Engine:**

- Domain-independent algorithms.
    - Derives new information:
    - From above KB → concludes “The ground is wet.”

- **Agent Program (Decision Maker):**

- Chooses actions based on knowledge + goals.
    - “If ground is wet → carry umbrella.”

# Generic Knowledge-based Agent



```
function KB-AGENT(percept) returns an action
  persistent: KB, a knowledge base
  t, a counter, initially 0, indicating time
  TELL(KB, MAKE-PERCEP-SENTENCE(percept, t))
  action  $\leftarrow$  ASK(KB, MAKE-ACTION-QUERY(t))
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))
  t  $\leftarrow$  t + 1
  return action
```

- Memorize percept at time *t*
- Ask for logical action given an objective
- Record action taken at time *t*

# Outline

## Introduction to Knowledge and Reasoning

### Formal Foundations of Knowledge Representation

- Facts, Models, Entailment
- Knowledge based Agents

### Knowledge Representation Techniques

- Logic (Prepositional & First Order)
- Semantic Networks, Frames, Rules, Ontologies

### Reasoning Approaches

- Deductive, Inductive, Abductive
- Inference Procedures
- Soundness & Completeness

### Modern Views of KR

- Large Language Models (LLMs)
- Probabilistic Reasoning

# Logic (Symbolic / Declarative Representation)

- Logic is a **formal system for representing knowledge** in a precise, unambiguous, and mathematically structured way.
- It allows an intelligent agent to:
  - Represent facts and rules about the world, and
  - **Reason** by drawing valid conclusions from them.
- It is called **declarative** because you *declare* what is true, and the reasoning system figures out the consequences automatically.
- **Types of Logic in AI**
  - **Propositional Logic (PL)**
  - **First-Order Logic (FOL)**

# Propositional Logic (PL)

- **Simplest form of logic**; deals with *propositions* (statements that can be True or False).
- **Syntax**
  - **Atomic propositions** (e.g., P, Q, R).
  - **Logical connectives**  $\neg$  (NOT),  $\wedge$  (AND),  $\vee$  (OR),  $\rightarrow$  (IMPLIES),  $\leftrightarrow$  (IFF).
- **Semantics**
  - Each proposition has a truth value (T/F).
  - The truth of complex statements is determined by **truth tables**
    - $P$ : “*It is raining.*”
    - $Q$ : “*The ground is wet.*”
    - *Rule*:  $P \rightarrow Q$
    - *If P = True  $\rightarrow$  Q must also be True (Entailment).*
- *Interpretation*  $\square$  The agent “knows” that if it is raining, the ground becomes wet.

# First-Order Logic (FOL)

- Extends propositional logic by introducing **objects, relations, and quantifiers**.
- **Why needed?**
  - PL cannot represent general statements like “All humans are mortal.”
- **Components of FOL**
  - **Constants:** Represent objects (e.g., Socrates).
  - **Predicates:** Represent relationships/properties (e.g.,  $\text{Human}(x)$ ,  $\text{Mortal}(x)$ ).
  - **Quantifiers**
    - **Universal ( $\forall$ ):** For all
    - **Existential ( $\exists$ ):** There exists
      - $\forall x (\text{Human}(x) \rightarrow \text{Mortal}(x))$
      - $\text{Human}(\text{Socrates})$
      - $\models \text{Mortal}(\text{Socrates})$  (*Socrates is mortal*)
- **Interpretation**  We can describe general truths and infer facts about specific individuals.

# Logic (Symbolic / Declarative Representation) (cont..)

- **Reasoning in Logic**

- **Deduction:** Deriving specific facts from general rules.
- Applies logical rules to move from general knowledge to specific instances.
  - General Rule:  $\forall x (\text{Human}(x) \rightarrow \text{Mortal}(x))$  — *All humans are mortal.*
  - Fact:  $\text{Human}(\text{Socrates})$  — *Socrates is a human.*
  - Deduction:  $\text{Mortal}(\text{Socrates})$  — *Therefore, Socrates is mortal.*

- **Inference Rules:** Formal procedures used to draw conclusions

- **Modus Ponens:** If  $P \rightarrow Q$  and  $P$ , then  $Q$ .  
E.g. If it rains, the ground is wet  $\rightarrow$  It rains  $\rightarrow$  The ground is wet.
- **Resolution:** Combine clauses to infer new facts ( $\neg P \vee Q$  and  $P \Rightarrow Q$ ).
- **Unification:** Match variables with constants ( $x = \text{Socrates}$ ).
- **Entailment ( $\models$ ):** A sentence is entailed if it's true in all models of the KB.
  - $\text{KB} = \{\text{It rains} \rightarrow \text{ground wet}, \text{It rains}\} \models \text{ground wet}$
- Logical reasoning ensures that every inferred statement follows necessarily from the knowledge base.

# Logic (Symbolic / Declarative Representation) (cont..)

- **Strengths**

- Mathematically **exact**; statements have **clear** and **precise** meanings.
- Inference produces **correct** and **exhaustive conclusions**.
- Rules can be combined to infer **new knowledge automatically**.
- Forms the basis for **expert systems**, **theorem provers**, and **planning systems**.

- **Limitations**

- Real-world knowledge is rarely **fully true or false**, making systems **brittle**.
- Cannot represent **uncertainty** such as “probably” or “likely.”
- Inference processes can become **computationally expensive** as **knowledge size increases**.
- Requires experts to manually **encode facts**, making **knowledge acquisition** difficult.

- **Applications**

- **Expert systems** □ MYCIN for medical diagnosis.
- **Automated reasoning** □ theorem provers, SAT solvers.
- **Knowledge-based planning systems**.
- **Semantic web ontologies** □ foundation of OWL logic.

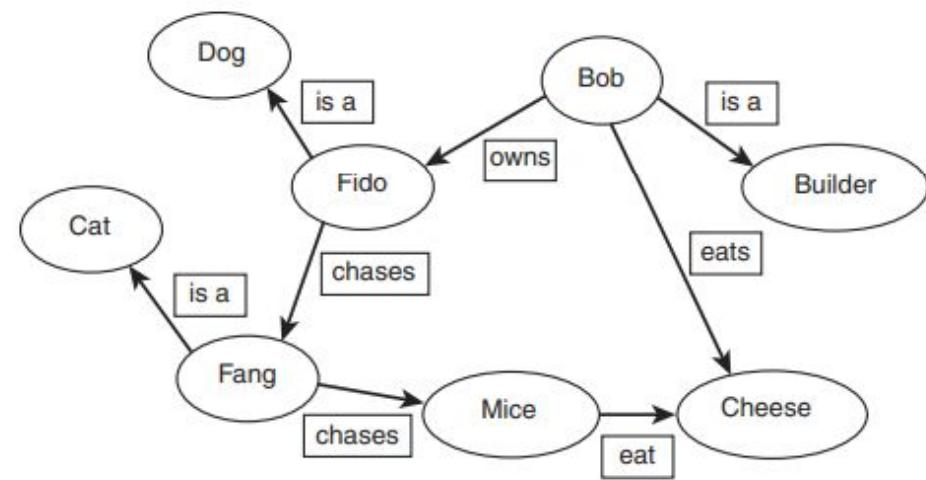
# Semantic Networks (Graph-Based Representation)

- A **Semantic Network** is a **graph-based knowledge representation** model where
  - **Nodes** represent **concepts or entities**, and
  - **Edges (links)** represent **relationships** between these concepts.
- It provides a **visual and structural way** to represent how knowledge is organized and related in the human mind or in an AI system.
- *Interpretation* □ Instead of storing knowledge as text or equations, we represent it as a **network of connected meanings** — hence the term “*semantic*” (meaning-based)
- **Structure of a Semantic Network**

Element	Description
<b>Node</b>	Object, concept, or class
<b>Edge / Link</b>	Relationship between nodes
<b>Inheritance</b>	Lower nodes inherit properties of higher nodes

# Semantic Networks (Graph-Based Representation) (cont..)

- This network shows how **objects and relationships** are represented in AI.
- **Nodes** represent **concepts or entities** (e.g., *Bob*, *Fido*, *Cheese*).
- **Links** show **relationships** (e.g., *is-a*, *owns*, *chases*, *eats*).
- Knowledge flows through these connections, allowing reasoning and inheritance.



# Semantic Networks (Graph-Based Representation) (cont..)

- **Reasoning in Semantic Networks**
- Semantic networks support **reasoning through inheritance** and **associative links**:
  - **Inheritance**: Lower-level nodes automatically acquire properties from higher-level nodes.
  - *Since Fido is a Dog, and Dogs are Animals, Fido inherits all animal properties (e.g., needs food, can move).*
  - *Fido → Dog → Animal → inherits properties like “can move” and “needs food.”*
  - **Spreading Activation**: When one concept is activated, related concepts are also activated.
  - *Activating Bob in the network may trigger related nodes like Fido (his pet), Builder (his occupation), or Cheese (something he eats).*

# Semantic Networks (Graph-Based Representation) (cont..)

- **Types of Semantic Networks**

- **Definitional Networks**

- Represent “is-a” hierarchies and taxonomies.
    - Dog  $\square$  Mammal  $\square$  Animal

- **Assertional Networks**

- Represent specific facts about instances.
    - Fido  $\square$  instance-of  $\square$  Dog

- **Inheritance Networks**

- Represent property inheritance and exceptions.
    - Bird  $\square$  can fly
    - Penguin  $\square$  Bird, but *cannot fly* (exception handling)

- **Conceptual Dependency Networks**

- Represent actions and relations (used in NLP).
    - “John eats an apple”  $\square$  (ACTOR: John, ACTION: Eat, OBJECT: Apple)

# Semantic Networks (Graph-Based Representation) (cont..)

- **Strengths**

- Intuitive and visual; easy to understand for both humans and machines.
- Supports inheritance by automatically transferring properties down hierarchy.
- Efficient for retrieval and quick association of related concepts.
- Useful in applications like semantic search and knowledge graphs.

- **Limitations**

- Relationships between nodes may lack clear or formal semantics.
- Cannot easily represent complex logical constraints.
- Does not handle uncertainty; all relationships are assumed definite.
- Scalability issues arise as large networks become dense and difficult to maintain.

- **Applications**

- **WordNet** □ Large lexical database linking words by meaning (used in NLP).
- **Google Knowledge Graph** □ Represents entities and their relations for semantic search.
- **Expert Systems** □ Conceptual understanding in diagnosis or classification.
- **ConceptNet** □ Open-source common-sense reasoning network.

# Frames (Structured Representation)

- Frames are **data structures** that represent **structured knowledge** about objects, situations, or concepts in terms of **attributes (called slots)** and their **corresponding values**.  
Each frame acts like a **template** or **schema** that describes an entity and its properties.
- It allows an intelligent system to:
  - Store and organize knowledge in a hierarchical way.
  - Use **default reasoning** and **inheritance** to fill missing details.
- *Interpretation* □ *Frames are similar to “objects” in object-oriented programming, where each object has attributes, default values, and possibly relationships to other objects.*
- **Structure of a Frame** A frame consists of
  - **Slots** □ Attributes or properties of the object.
  - **Values** □ Specific information or default values for each slot.
  - **Inheritance Links** □ Connect frames in a hierarchy, allowing lower-level frames to inherit properties from higher-level frames.

# Frames (Structured Representation) (cont..)

- The **Vehicle** frame defines general attributes (slots).
- The **Car** frame inherits all attributes from Vehicle, adds new ones (Type, Color), and can override defaults.
- The **ElectricCar** frame further inherits and overrides (Engine → Electric).
- When a query is made for a property, the system uses:
  - **Instance value**, if available.
  - **Default value**, if not specified.
  - **Inherited value**, if neither of the above is defined.

## Vehicle (Frame)

- Wheels: 4
- Engine: Generic
- HasDoors: Yes



## Car (Frame)

- Inherits from: Vehicle
- Engine: Petrol
- Type: Sedan
- Color: White



## Electric Car (Frame)

- Inherits from: Car
- Engine: Electric (Overrides)
- Battery: Lithium-Ion
- Range: 350 km

# Frames (Structured Representation) (cont..)

- **Reasoning with Frames**
- Reasoning in frames allows AI systems to **reuse, extend, and infer knowledge** efficiently through inheritance and default logic.
  - **Inheritance:** Lower-level frames automatically acquire attributes from higher-level frames.
  - The **Car** frame inherits Wheels: 4 and HasDoors: Yes from **Vehicle**.
  - **Default Reasoning:** If a slot has no specific value, the default is used.
  - The **Car** frame uses Color: White unless another value is specified.
- **Procedural Attachment:** Frames can trigger actions when a slot is accessed or updated.
- Accessing the **Engine** slot in **ElectricCar** may call a procedure to compute Battery Range.

# Frames (Structured Representation) (cont..)

- **Strengths**

- Organizes knowledge in a **structured** and **hierarchical** way.
- Supports **default reasoning** and **inheritance**.
- Captures **common-sense** and **contextual knowledge**.
- Easy to **understand** and **maintain** due to object-like structure.

- **Limitations**

- Difficult to handle **uncertainty** or **changing information**.
- Can become **complex** in large hierarchies.
- Not ideal for representing **procedural** or **relational knowledge** (e.g., time-based or logical relationships).

- **Applications**

- **Expert systems** □ Medical or engineering diagnosis.
- **Computer vision** □ Scene and object understanding.
- **Natural language processing** □ Context and meaning representation.
- **Knowledge-based robotics** □ Representing structured world models.

# Production Rules (Rule-Based Systems)

- Production rules represent **procedural knowledge** — what actions to take when specific conditions are met.
- They are expressed as **IF–THEN statements**, describing a cause–effect or condition–action relationship.
- *Interpretation* □ Rules simulate human decision-making logic and are widely used in expert systems.
- **Structure of a Rule**
  - A production rule has two parts:
  - **Condition (IF)** □ The situation or fact to be checked.
  - **Action (THEN)** □ The conclusion or response triggered when the condition is satisfied.

# Production Rules (Rule-Based Systems) (cont..)

- **Rule 1:** IF fever AND cough THEN possible flu
- **Rule 2:** IF temperature > 38°C THEN fever = True
- **Known Facts**
  - temperature = 39°C
  - cough = True
- **Reasoning Flow**
  - The system checks Rule 2 → since temperature > 38°C → infer **fever = True**.
  - Now Rule 1 applies → since **fever = True** and **cough = True** → infer **possible flu**.

# Production Rules (Rule-Based Systems) (cont..)

- **Reasoning in Rule-Based Systems**

- There are two main reasoning strategies:
- **Forward Chaining (Data-Driven):** Starts with **known facts** and applies rules to derive new facts.
  - Given  $temperature > 38^{\circ}\text{C}$ , infer  $fever = \text{True}$ , then infer  $possible\ flu$ .
- **Backward Chaining (Goal-Driven):** Starts with a **goal** and works **backward** to determine if facts support it.
  - To test “Does the patient have flu?”, the system checks rules that conclude flu.

# Production Rules (Rule-Based Systems) (cont..)

- **Strengths**

- Simple and **intuitive** to design (“if–then” format).
- Supports **step-by-step reasoning**.
- Easy to **modify** or add new rules.
- **Transparent reasoning process** — easy to explain conclusions.

- **Limitations**

- Large systems can contain **thousands of rules** that are hard to manage.
- Multiple rules may **trigger simultaneously**.
- Cannot easily handle **probabilistic or uncertain reasoning**.

- **Applications**

- **Expert systems** □ MYCIN (medical diagnosis), DENDRAL (chemical analysis).
- **Decision support systems** □ Business and industrial rule engines.
- **Control systems** □ Troubleshooting, maintenance, configuration.
- **Game AI** □ Strategy-based action triggering.

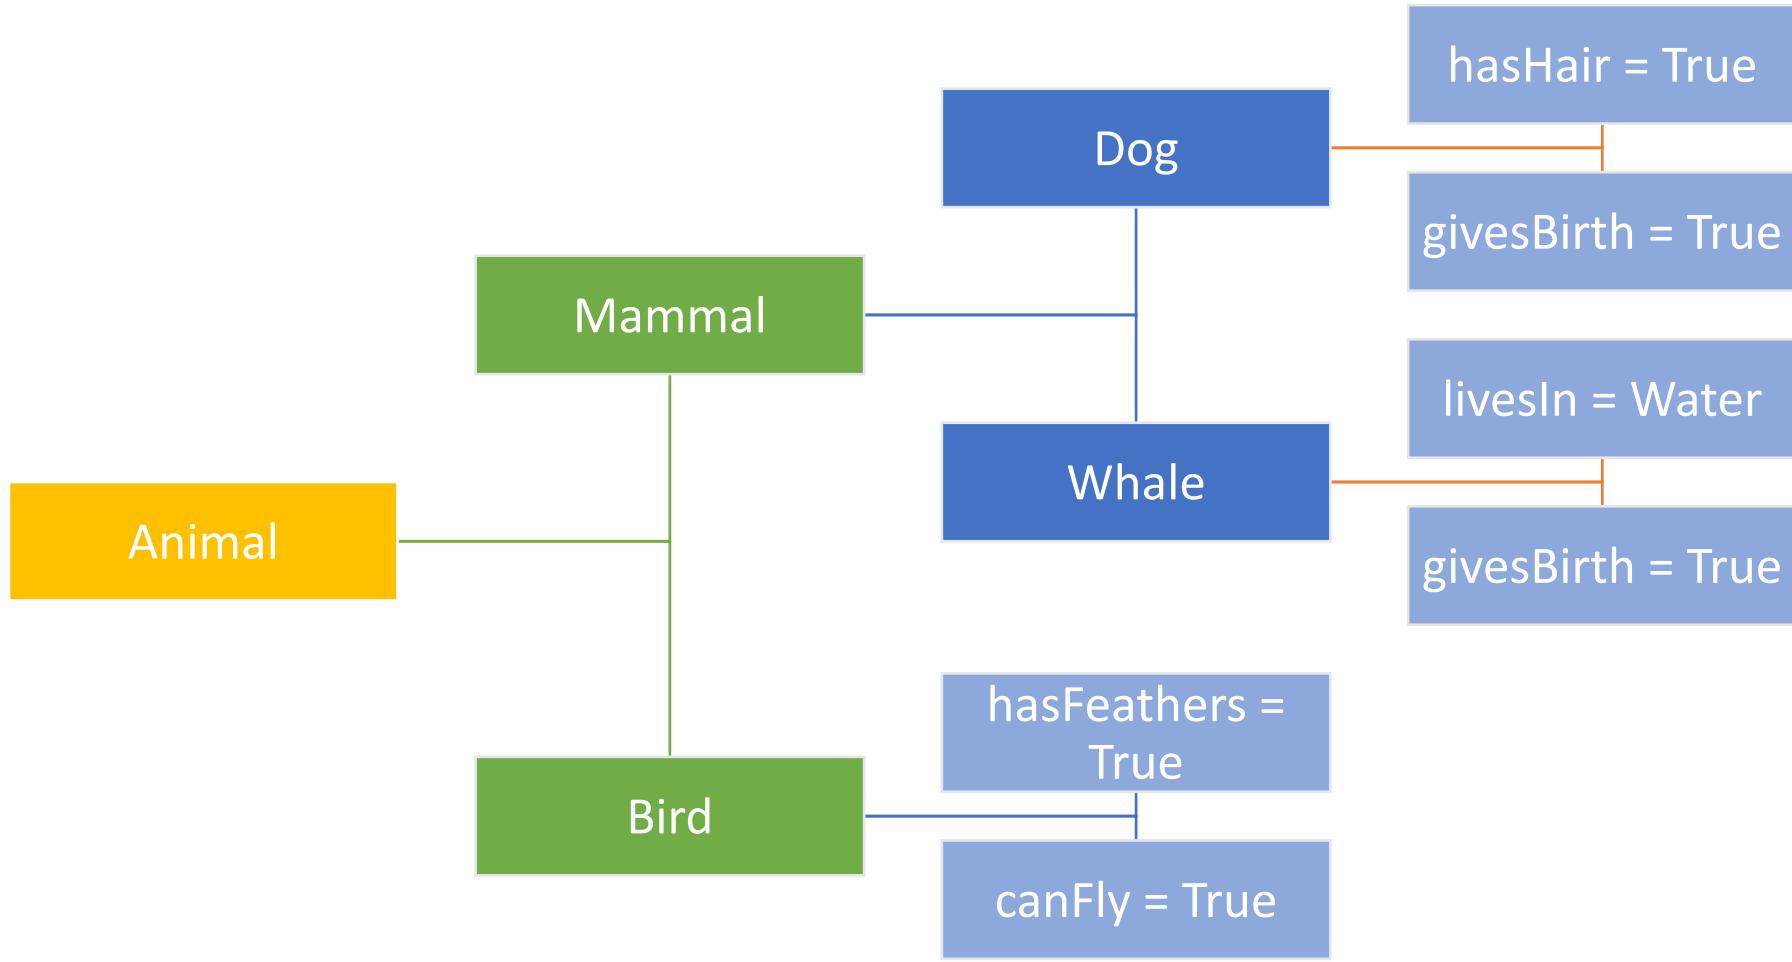
# Ontologies (Hierarchical Representation)

- An **ontology** is a **formal, explicit, and hierarchical representation** of knowledge within a domain. It defines **concepts (classes), relationships, properties, and instances** to enable machines to understand and reason about domain information.
- *Interpretation* □ Ontologies provide a **shared vocabulary** and **structure** that allows knowledge to be reused and exchanged between systems.
- **Components of an Ontology**

Component	Description
Classes	Abstract categories of objects
Subclasses	More specific classes within a category
Relations	Define how classes connect
Properties	Attributes or characteristics
Instances	Concrete examples of a class

# Ontologies (Hierarchical Representation) (cont..)

- Dog and Whale inherit properties from Mammal and Animal.
- Ontology allows reasoning such as:
  - If an entity is a Dog, it is a Mammal and therefore an Animal.
  - If Mammals give birth, we can infer that Dogs give birth too.



# Ontologies (Hierarchical Representation) (cont..)

- **Reasoning with Ontologies**
- Enables **logical inference**, **error detection**, and **automatic classification** across related concepts.
  - **Inheritance:** Subclasses inherit properties from parent classes.
  - Since “Dog” is a subclass of “Mammal,” it inherits givesBirth = True and hasHair = True.
  - **Classification:** Determines where new instances belong in the hierarchy.
  - If “Tweety” has hasFeathers = True and canFly = True, it is classified as a Bird.
- **Consistency Checking:** Ensures that relationships don’t contradict each other.
- An instance cannot be both a Bird and a Fish if those classes are defined as disjoint.
- **Semantic Inference:** Discovers new knowledge from existing relationships.
- If “All mammals are animals” and “Dogs are mammals,” the system infers that Dogs are animals.

# Ontologies (Hierarchical Representation) (cont..)

- **Strengths**

- Provides **shared** and reusable domain knowledge.
- Enables **semantic interoperability** between systems.
- Supports **reasoning**, **classification**, and **data integration**.
- Essential for **Semantic Web** and **NLP**.

- **Limitations**

- Requires significant **domain expertise** to design.
- **Complex** and **time-consuming** to maintain and update.
- **Limited scalability** for very large or dynamic domains.

- **Applications**

- Semantic Web technologies □ RDF, OWL, SPARQL.
- Healthcare ontologies □ SNOMED CT, Gene Ontology.
- Knowledge graphs □ Google Knowledge Graph, DBpedia.
- Natural Language Processing □ Concept linking, entity recognition.

# Outline

## Introduction to Knowledge and Reasoning

### Formal Foundations of Knowledge Representation

- Facts, Models, Entailment
- Knowledge based Agents

### Knowledge Representation Techniques

- Logic (Prepositional & First Order)
- Semantic Networks, Frames, Rules, Ontologies

### Reasoning Approaches

- Deductive, Inductive, Abductive
- Inference Procedures
- Soundness & Completeness

### Modern Views of KR

- Large Language Models (LLMs)
- Probabilistic Reasoning

# Reasoning Approaches in AI

- Reasoning is the process through which an intelligent system **derives conclusions** or **makes decisions** based on existing knowledge and evidence.
- Different reasoning approaches define how conclusions are reached — from general to specific, or from effects to causes.
- **Types of Reasoning**
  - **Deductive Reasoning** (Top-Down)
    - Draws **specific conclusions** from **general rules**.
    - **Form:** If  $A \rightarrow B$ , and  $A$  is true, then  $B$  is true.
    - **Rule:** All birds can fly.  
**Fact:** Parrot is a bird.  
⇒ Conclusion: Parrot can fly.
    - *Used for logical, certain conclusions.*

# Reasoning Approaches in AI (cont..)

- **Inductive Reasoning** (Bottom-Up)
  - Generalizes **rules or patterns** from **observed examples or data**.
  - **Observation**: Swan<sub>1</sub>, Swan<sub>2</sub>, Swan<sub>3</sub> are white.  
⇒ **Conclusion**: All swans are white.
  - *Used in machine learning — patterns inferred from data.*
- **Abductive Reasoning** (Best Explanation)
  - Infers the **most likely cause** from an observation.
  - **Observation**: Ground is wet.  
**Possible causes**: It rained, sprinkler was on, someone spilled water.  
⇒ **Most likely**: It rained.
  - *Used in diagnostic systems and hypothesis generation.*

# Reasoning Approaches in AI (cont..)

- **Inference Procedures**

- Forward Chaining (Data-Driven)

- Starts with **known facts** → applies **rules** → derives **new conclusions**.

- Used in production systems (e.g., MYCIN).

- Backward Chaining (Goal-Driven)

- Starts with a **goal** → works **backward** → checks **supporting facts**.

- Used in systems like Prolog.

- Resolution

- A rule-based method to **prove logical conclusions** using **contradictions**.

- Used in automated theorem proving.

# Reasoning Approaches in AI (cont..)

- **Soundness & Completeness**
- **Soundness**
  - If the inference procedure derives a conclusion, that conclusion is **logically correct** (no false results).
- **Completeness**
  - If a conclusion is **logically entailed**, the inference procedure **can derive it** (no missing truths).
- *Goal in AI reasoning: inference systems should be both **sound** and **complete**.*

# Outline

## Introduction to Knowledge and Reasoning

### Formal Foundations of Knowledge Representation

- Facts, Models, Entailment
- Knowledge based Agents

### Knowledge Representation Techniques

- Logic (Prepositional & First Order)
- Semantic Networks, Frames, Rules, Ontologies

### Reasoning Approaches

- Deductive, Inductive, Abductive
- Inference Procedures
- Soundness & Completeness

### Modern Views of KR

- Large Language Models (LLMs)
- Probabilistic Reasoning

# Modern Views of KR

- Modern approaches to Knowledge Representation (KR) go beyond symbolic systems like logic and semantic networks.
- Integrate **probabilistic models**, **neural networks**, and **Large Language Models (LLMs)**.
- Enable reasoning under **uncertainty** and **incomplete information**.
- Provide **scalability** to handle large, complex real-world data.
- Learn and represent knowledge **automatically from data** rather than predefined rules.
- Bridge the gap between **symbolic reasoning** and **data-driven learning**.
- Traditional KR focused on explicit rules; modern KR emphasizes **learning representations automatically from data**.

# Large Language Models (LLMs)

- LLMs (e.g., GPT, BERT, PaLM) represent knowledge implicitly within **high-dimensional vector spaces** learned from massive text data.
- They encode **semantics**, **context**, and **relationships** without needing predefined rules.
- Learn from **large-scale data** (trillions of tokens).
- Represent meaning through **embeddings** (numerical representations of words/concepts).
- Enable reasoning via **contextual understanding** and **in-context learning**.
- **Prompt:** “If Paris is in France, where is Berlin?”  LLM infers: “Germany” — by associating learned patterns rather than explicit logic rules.

# Large Language Models (LLMs) (cont..)

- **Advantages**

- Capture **real-world knowledge** implicitly.
- Handle **ambiguity** and **natural language** variability.
- Continuously **adaptable** to new data.

- **Limitations**

- Lack explicit **reasoning** transparency.
  - May generate plausible but **incorrect answers** (hallucination).
  - Knowledge not easily **inspectable** or **modifiable**.
- *Interpretation* □ LLMs shift KR from symbolic **representation** to **statistical and neural representation**, bridging knowledge and language understanding.

# Probabilistic Reasoning

- Traditional logic deals with *certainty* (True/False).
- **Probabilistic reasoning** handles *uncertainty* using mathematical probability to infer likely conclusions.

- **Core Concept**

Instead of “The patient has flu,” we represent:

$$P(\text{Flu} \mid \text{Fever, Cough}) = 0.85$$

- **Common Models**

- **Bayesian Networks**  Graphical models representing probabilistic dependencies.
- **Hidden Markov Models (HMMs)**  Used for sequential reasoning (e.g., speech recognition).
- **Markov Decision Processes (MDPs)**  Used in AI planning and reinforcement learning.

# Probabilistic Reasoning (cont..)

- **Example (Bayesian Inference)**

- If 90% of flu patients have fever, and 10% of non-flu patients have fever, then observing “fever” increases belief in “flu” but doesn’t guarantee it.

- **Advantages**

- Handles **noisy** or **incomplete** data.
- Supports **decision-making** under uncertainty.
- Widely used in real-world AI (diagnosis, prediction, robotics).

- **Limitations**

- Requires **large**, accurate **probability** data.
  - Computationally **intensive** for large networks.
- *Interpretation*  Probabilistic reasoning combines logic and statistics to make **rational decisions** under uncertainty — a foundation for **modern intelligent systems**.

# Integration of Symbolic + Neural KR (Neuro-Symbolic AI)

- Modern AI research seeks to **combine structured symbolic reasoning** (logic, ontologies) with **data-driven learning** (LLMs, neural nets).
- **Goal**
  - Achieve systems that can **learn, reason, and explain**.
- Use an LLM for text understanding and a logic layer for ensuring factual consistency.
- Neuro-symbolic systems represent the **future of KR**, enabling AI that is both **intelligent and interpretable**.

# Summary

- We explored
  - **Knowledge & Reasoning:** How systems represent, reason, and derive conclusions using facts, models, and entailment.
  - **Representation Techniques:** Logic, Semantic Networks, Frames, Rules, and Ontologies for structured and relational knowledge.
  - **Reasoning Approaches:** Deductive, Inductive, Abductive reasoning; Forward/Backward Chaining; Soundness & Completeness.
  - **Modern KR:** Shift to data-driven models — LLMs for language-based reasoning and Probabilistic methods for uncertainty handling.
  - **Trend:** Integration of **symbolic** and **neural** reasoning — forming the basis of **Neuro-Symbolic AI**.
  - *“Traditional AI focused on representing what we know; modern AI learns how to represent knowledge itself — bridging logic, probability, and language to achieve true understanding.”*